

```

from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

!pip install pydub python_speech_features

Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting pydub
  Downloading pydub-0.25.1-py2.py3-none-any.whl (32 kB)
Collecting python_speech_features
  Downloading python_speech_features-0.6.tar.gz (5.6 kB)
  Preparing metadata (setup.py) ... done
Building wheels for collected packages: python_speech_features
  Building wheel for python_speech_features (setup.py) ... done
  Created wheel for python_speech_features: filename=python_speech_features-0.6-py3-none-any.whl size=5886 sha256=2726ece95ff901b839d16
  Stored in directory: /root/.cache/pip/wheels/09/a1/04/08e2688d2562d8f9ff89e77c6ddfbf7268e07dae1a6f22455e
Successfully built python_speech_features
Installing collected packages: python_speech_features, pydub
Successfully installed pydub-0.25.1 python_speech_features-0.6

```

▼ Importing Packages

```

import os
import pydub
from pydub import AudioSegment
from pydub.silence import split_on_silence
from python_speech_features import mfcc
import matplotlib.pyplot as plt
import numpy as np
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
from keras.utils import to_categorical
from sklearn.metrics import classification_report
import pandas as pd
import tensorflow as tf
from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation
from sklearn.model_selection import train_test_split
from keras.models import load_model
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import confusion_matrix
import warnings
warnings.filterwarnings('ignore')

```

▼ Data Preprocessing

```

# Set directory path
directory = '/content/drive/MyDrive/audio/'

# Get list of all folders in directory
class_labels = [folder for folder in os.listdir(directory) if os.path.isdir(os.path.join(directory, folder))]

# Print list of folders
print(class_labels)

['.ipynb_checkpoints', 'Siren', 'Dog_bark', 'Car_honk']

# Set parameters
n_mfcc = 20

# Initialize lists for audio data and labels
X = []
y = []

# Loop through each class directory
for i, label in enumerate(class_labels):
    class_dir = f'/content/drive/MyDrive/audio/{label}/'

```

```
# Loop through each audio file in the class directory
for filename in os.listdir(class_dir):
    file_path = os.path.join(class_dir, filename)

    try:
        # Load audio file
        audio = AudioSegment.from_file(file_path)
        audio_data = np.array(audio.get_array_of_samples())
        sample_rate = audio.frame_rate

        # Split audio into chunks of silence
        audio_chunks = split_on_silence(audio, min_silence_len=500, silence_thresh=-50)

        # Extract MFCCs from each chunk of audio
        nfft = 2400 # Set larger FFT size
        for chunk in audio_chunks:
            chunk_data = np.array(chunk.get_array_of_samples())
            mfccs = mfcc(chunk_data, samplerate=sample_rate, numcep=n_mfcc, nfft=nfft) # Set larger nfft value
            mfccs = np.mean(mfccs, axis=0)

            # Append MFCCs and label to lists
            X.append(mfccs)
            y.append(i)

    except Exception as e:
        # Skip file if it is not a recognized audio format
        print(f'Skipping file {filename}: {e}')

Skipping file 421017__girlwithsoundrecorder__ambulance-in-oradea.wav: 64
WARNING:root:frame length (4800) is greater than FFT size (2400), frame will be truncated. Increase NFFT to avoid.
WARNING:root:frame length (4800) is greater than FFT size (2400), frame will be truncated. Increase NFFT to avoid.
WARNING:root:frame length (4800) is greater than FFT size (2400), frame will be truncated. Increase NFFT to avoid.
WARNING:root:frame length (4800) is greater than FFT size (2400), frame will be truncated. Increase NFFT to avoid.
```

X

y

```
# Convert labels to one-hot encoded vectors
label_encoder = LabelEncoder()
y = label_encoder.fit_transform(y)

onehot_encoder = OneHotEncoder(sparse=False)
y = onehot_encoder.fit_transform(y.reshape(len(y), 1))
```

```

y
array([[1., 0., 0.],
       [1., 0., 0.],
       [1., 0., 0.],
       ...,
       [0., 0., 1.],
       [0., 0., 1.],
       [0., 0., 1.]])

# Convert lists to numpy arrays
X = np.array(X)
y = np.array(y)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

X_train.shape

(630, 20)

```

▼ Model Training and Testing

```

# Define the model architecture
model = Sequential()
model.add(Dense(256, input_shape=(X.shape[1],)))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(128))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(3))
model.add(Activation('softmax'))

# Compile the model
model.compile(loss='categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 256)	5376
activation (Activation)	(None, 256)	0
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 128)	32896
activation_1 (Activation)	(None, 128)	0
dropout_1 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 3)	387
activation_2 (Activation)	(None, 3)	0
=====		
Total params: 38,659		
Trainable params: 38,659		
Non-trainable params: 0		

```

# Train the model
history = model.fit(X_train, y_train,
                    batch_size=32,
                    epochs=100,
                    validation_data=(X_test, y_test))

```

```

Epoch 73/100
20/20 [=====] - 0s 6ms/step - loss: 0.2637 - accuracy: 0.8968 - val_loss: 0.4316 - val_accuracy: 0.8963
Epoch 74/100
20/20 [=====] - 0s 5ms/step - loss: 0.2829 - accuracy: 0.8921 - val_loss: 0.4284 - val_accuracy: 0.9037
Epoch 75/100
20/20 [=====] - 0s 5ms/step - loss: 0.2602 - accuracy: 0.9143 - val_loss: 0.4333 - val_accuracy: 0.8963
Epoch 76/100
20/20 [=====] - 0s 6ms/step - loss: 0.2463 - accuracy: 0.9143 - val_loss: 0.4336 - val_accuracy: 0.9074
Epoch 77/100
20/20 [=====] - 0s 6ms/step - loss: 0.2669 - accuracy: 0.9032 - val_loss: 0.4177 - val_accuracy: 0.9037
Epoch 78/100
20/20 [=====] - 0s 5ms/step - loss: 0.2523 - accuracy: 0.9016 - val_loss: 0.4244 - val_accuracy: 0.9000
Epoch 79/100
20/20 [=====] - 0s 6ms/step - loss: 0.2485 - accuracy: 0.9079 - val_loss: 0.4455 - val_accuracy: 0.8963
Epoch 80/100
20/20 [=====] - 0s 6ms/step - loss: 0.2682 - accuracy: 0.9032 - val_loss: 0.4290 - val_accuracy: 0.8926
Epoch 81/100
20/20 [=====] - 0s 6ms/step - loss: 0.2400 - accuracy: 0.9143 - val_loss: 0.4228 - val_accuracy: 0.9074
Epoch 82/100
20/20 [=====] - 0s 5ms/step - loss: 0.2540 - accuracy: 0.8984 - val_loss: 0.4333 - val_accuracy: 0.9037
Epoch 83/100
20/20 [=====] - 0s 5ms/step - loss: 0.2543 - accuracy: 0.9095 - val_loss: 0.4195 - val_accuracy: 0.9037
Epoch 84/100
20/20 [=====] - 0s 5ms/step - loss: 0.2262 - accuracy: 0.9159 - val_loss: 0.4368 - val_accuracy: 0.9000
Epoch 85/100
20/20 [=====] - 0s 6ms/step - loss: 0.2622 - accuracy: 0.8921 - val_loss: 0.4504 - val_accuracy: 0.9148
Epoch 86/100
20/20 [=====] - 0s 5ms/step - loss: 0.2582 - accuracy: 0.8905 - val_loss: 0.4414 - val_accuracy: 0.9000
Epoch 87/100
20/20 [=====] - 0s 5ms/step - loss: 0.2619 - accuracy: 0.9016 - val_loss: 0.4269 - val_accuracy: 0.9000
Epoch 88/100
20/20 [=====] - 0s 6ms/step - loss: 0.2483 - accuracy: 0.9095 - val_loss: 0.4376 - val_accuracy: 0.8852
Epoch 89/100
20/20 [=====] - 0s 5ms/step - loss: 0.2422 - accuracy: 0.9063 - val_loss: 0.4501 - val_accuracy: 0.8963
Epoch 90/100
20/20 [=====] - 0s 7ms/step - loss: 0.2212 - accuracy: 0.9032 - val_loss: 0.4351 - val_accuracy: 0.8963
Epoch 91/100
20/20 [=====] - 0s 10ms/step - loss: 0.2080 - accuracy: 0.9206 - val_loss: 0.4121 - val_accuracy: 0.9000
Epoch 92/100
20/20 [=====] - 0s 8ms/step - loss: 0.1964 - accuracy: 0.9159 - val_loss: 0.4236 - val_accuracy: 0.9111
Epoch 93/100
20/20 [=====] - 0s 8ms/step - loss: 0.2191 - accuracy: 0.9190 - val_loss: 0.4497 - val_accuracy: 0.9111
Epoch 94/100
20/20 [=====] - 0s 9ms/step - loss: 0.1915 - accuracy: 0.9302 - val_loss: 0.4726 - val_accuracy: 0.9074
Epoch 95/100
20/20 [=====] - 0s 7ms/step - loss: 0.2401 - accuracy: 0.9127 - val_loss: 0.4470 - val_accuracy: 0.9000
Epoch 96/100
20/20 [=====] - 0s 8ms/step - loss: 0.2035 - accuracy: 0.9238 - val_loss: 0.4640 - val_accuracy: 0.9074
Epoch 97/100
20/20 [=====] - 0s 9ms/step - loss: 0.2249 - accuracy: 0.9016 - val_loss: 0.4682 - val_accuracy: 0.8889
Epoch 98/100
20/20 [=====] - 0s 10ms/step - loss: 0.2461 - accuracy: 0.9111 - val_loss: 0.4774 - val_accuracy: 0.9074
Epoch 99/100
20/20 [=====] - 0s 8ms/step - loss: 0.2390 - accuracy: 0.9079 - val_loss: 0.4781 - val_accuracy: 0.8926
Epoch 100/100
20/20 [=====] - 0s 9ms/step - loss: 0.1935 - accuracy: 0.9238 - val_loss: 0.4689 - val_accuracy: 0.9037

```

```


history_df = pd.DataFrame(history.history)
history_df[['accuracy', 'val_accuracy']].plot()

```

```
<Axes: >
```

```
history_df = pd.DataFrame(history.history)
history_df[['loss', 'val_loss']].plot()
```

```
<Axes: >
```



```
# Evaluate the model on the test set
test_loss, test_acc = model.evaluate(X_test, y_test)
print('Test accuracy:', test_acc)
```

```
9/9 [=====] - 0s 2ms/step - loss: 0.4689 - accuracy: 0.9037
Test accuracy: 0.9037036895751953
```

```
# Generate predictions on the test set
y_pred = model.predict(X_test)
```

```
# Convert predictions to class labels
y_pred_classes = np.argmax(y_pred, axis=1)
```

```
# Convert predictions to class labels
y_test_classes = np.argmax(y_test, axis=1)
```

```
# Print the classification report
print(classification_report(y_test_classes, y_pred_classes, target_names=['Siren', 'Dog_bark', 'Car_honk']))
```

```
9/9 [=====] - 0s 2ms/step
```

	precision	recall	f1-score	support
Siren	0.86	0.71	0.78	42
Dog_bark	0.93	0.97	0.95	171
Car_honk	0.86	0.84	0.85	57
accuracy			0.90	270
macro avg	0.88	0.84	0.86	270
weighted avg	0.90	0.90	0.90	270

```
class_labels = ['Siren', 'Dog_bark', 'Car_honk']
# Compute the confusion matrix
conf_matrix = confusion_matrix(y_test_classes, y_pred_classes)
```

```
# Plot the confusion matrix
fig, ax = plt.subplots(figsize=(5, 5))
im = ax.imshow(conf_matrix, cmap="Blues")
```

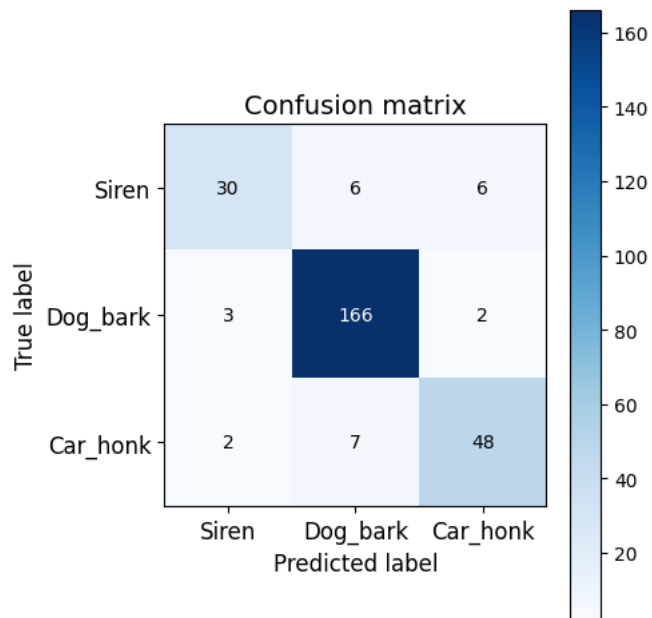
```
# Add colorbar
cbar = ax.figure.colorbar(im, ax=ax)
```

```
# Add labels to the x-axis and y-axis
ax.set_xticks(np.arange(len(class_labels)))
ax.set_yticks(np.arange(len(class_labels)))
ax.set_xticklabels(class_labels, fontsize=12)
ax.set_yticklabels(class_labels, fontsize=12)
```

```
# Add annotations
thresh = conf_matrix.max() / 2.0
for i in range(len(class_labels)):
    for j in range(len(class_labels)):
        ax.text(j, i, format(conf_matrix[i, j], "d"),
                ha="center", va="center",
                color="white" if conf_matrix[i, j] > thresh else "black")

# Add title and axis labels
ax.set_title("Confusion matrix", fontsize=14)
ax.set_xlabel("Predicted label", fontsize=12)
ax.set_ylabel("True label", fontsize=12)

plt.tight_layout()
plt.show()
```



```
from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import label_binarize

# Binarize the true labels
y_test_binarized = label_binarize(y_test, classes=[0, 1, 2])

classes = class_labels

# Compute the predicted class probabilities for each sample
y_pred_prob = model.predict(X_test)

# Compute the ROC curve and AUC for each class
fpr = dict()
tpr = dict()
roc_auc = dict()
for i in range(len(classes)):
    fpr[i], tpr[i], _ = roc_curve(y_test_binarized[:, i], y_pred_prob[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# Compute micro-average ROC curve and AUC
fpr_micro, tpr_micro, _ = roc_curve(y_test_binarized.ravel(), y_pred_prob.ravel())
roc_auc_micro = auc(fpr_micro, tpr_micro)

# Compute macro-average ROC curve and AUC
fpr_macro = np.unique(np.concatenate([fpr[i] for i in range(len(classes))]))
tpr_macro = np.zeros_like(fpr_macro)
for i in range(len(classes)):
    tpr_macro += np.interp(fpr_macro, fpr[i], tpr[i])
tpr_macro /= len(classes)
roc_auc_macro = auc(fpr_macro, tpr_macro)

# Plot the ROC curves for each class
plt.figure()
```

```

colors = ['blue', 'red', 'green']
for i, color in zip(range(len(classes)), colors):
    plt.plot(fpr[i], tpr[i], color=color, lw=2,
             label='ROC curve of class {0} (area = {1:0.2f})'
             ''.format(classes[i], roc_auc[i]))


# Plot the micro-average ROC curve
plt.plot(fpr_micro, tpr_micro, color='deeppink', lw=2, linestyle='--',
         label='micro-average ROC curve (area = {0:0.2f})'
         ''.format(roc_auc_micro))

# Plot the macro-average ROC curve
plt.plot(fpr_macro, tpr_macro, color='navy', lw=2, linestyle='--',
         label='macro-average ROC curve (area = {0:0.2f})'
         ''.format(roc_auc_macro))

# Add labels and legend
plt.plot([0, 1], [0, 1], 'k--', lw=2)
plt.xlim([-0.05, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Multi-Class Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()

```

9/9 [=====] - 0s 4ms/step



```
model.save('model.h5')
```

▼ Prediction Function

```

def predict_class(audio_path):
    n_mfcc = 20

    # Load audio file
    audio = AudioSegment.from_file(audio_path)
    audio_data = np.array(audio.get_array_of_samples())
    sample_rate = audio.frame_rate

    # Split audio into chunks of silence
    audio_chunks = split_on_silence(audio, min_silence_len=500, silence_thresh=-50)

    # Extract MFCCs from each chunk of audio
    mfccs_list = []
    nfft = 2048 # Set larger FFT size
    for chunk in audio_chunks:
        chunk_data = np.array(chunk.get_array_of_samples())

```



```

mfccs = mfcc(chunk_data, samplerate=sample_rate, numcep=n_mfcc, nfft=nfft) # Set larger nfft value
mfccs = np.mean(mfccs, axis=0)
mfccs_list.append(mfccs)

# Convert MFCCs to a numpy array
mfccs_array = np.array(mfccs_list)

# Reshape MFCCs to have 1 channel
mfccs_array = mfccs_array.reshape(mfccs_array.shape[0], mfccs_array.shape[1], 1)

# Load trained model
model = load_model('model.h5')

# Get predicted probabilities for each class
probs = model.predict(mfccs_array)

# Convert probabilities to predicted class labels
predicted_classes = np.argmax(probs, axis=1)

# Convert predicted class labels to class names
class_names = ['Siren', 'Dog_bark', 'Car_honk']
predicted_class_names = [class_names[i] for i in predicted_classes]

return predicted_class_names

predict_class('/content/drive/MyDrive/audio/Car_honk/235506__ceberation__car-horn.wav')

1/1 [=====] - 0s 65ms/step
['Car_honk']

predict_class('/content/drive/MyDrive/audio/Dog_bark/100032__nfræe__rose_bark.wav')

1/1 [=====] - 0s 72ms/step
['Dog_bark']

predict_class('/content/drive/MyDrive/audio/Siren/108192__paubg_pou__ambulancia.wav')

1/1 [=====] - 0s 72ms/step
['Siren']

```

✓ 2s completed at 10:30 PM

